

EE360C: Algorithms

Priority Queues

Spring 2018

Department of Electrical and Computer Engineering
University of Texas at Austin

Motivation

Motivation: Stable Marriage

The stable marriage algorithm needs a data structure that maintains the dynamically changing set of all free men. The algorithm needs to be able to:

- add elements to the set
- delete elements from the set
- select an element from the set, based on some assigned *priority*

Motivation: Sort a List of Numbers

Sort

Instance: Nonempty list x_1, x_2, \dots, x_n of integers

Solution: A permutation y_1, y_2, \dots, y_n of x_1, x_2, \dots, x_n such that $y_i \leq y_{i+1}$ for all $1 \leq i < n$

Possible Algorithm

- Store all of the numbers in a data structure D
- Repeatedly find the smallest number in D , output it, and remove it

To get $O(n \log n)$ running time, each “find minimum” step must take $O(\log n)$ time

Candidate Data Structures for Sorting

The data structure we select must support inserting a new element, finding the minimum element, and deleting the minimum element.

- **List:** Insertion and deletion take $O(1)$ time, but finding the minimum requires scanning the list and takes $\Omega(n)$ time
- **Sorted array:** Finding the minimum takes $O(1)$ time, but insertion and deletion take $\Omega(n)$ time in the worst case

Priority Queue

Enter the Priority Queue

- Store a set S of elements, where each element v has a priority value $\text{key}(v)$
- Smaller key values denote higher priorities
- Operations supported:
 - find the element with the smallest key
 - remove the element with the smallest key
 - insert a new element
 - delete an element
- Key update and element deletion require knowledge of the position of the element in the priority queue

An Example Application

Consider the problem of real-time scheduling of processes on a computer

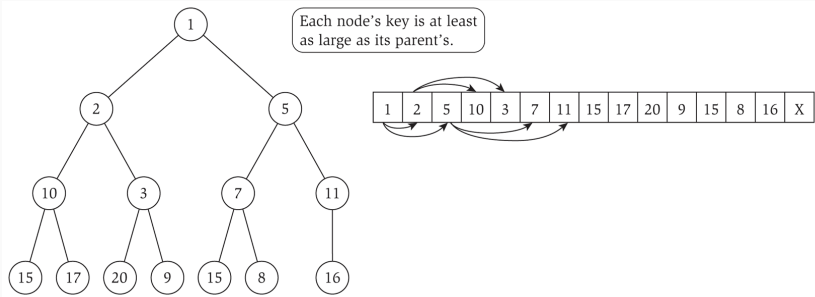
- each process has a priority
- processes *do not* arrive in order of their priorities
- we need to maintain a set of *active processes* with the ability to quickly extract the one with the highest priority so it can be scheduled
- using a priority queue keyed by process priority, scheduling the highest priority process entails simply finding the one with the lowest priority key

Heaps

Heaps

- Combine the benefits of both lists and sorted arrays
- Conceptually, a heap is a balanced binary tree
- **Heap order:** For every element v at node i , the element w at i 's parent satisfies $\text{key}(w) \leq \text{key}(v)$
- We can implement a heap in a pointer-based data structure
- Alternatively, assume a maximum number N of elements is known in advance
- Store nodes of the heap in an array
 - Node at index i has children at indices $2i$ and $2i + 1$ and parent at index $\lfloor i/2 \rfloor$
 - Index 1 is the root
 - How do you know that a node at index i is a leaf? If $2i > n$, the number of elements in the heap.

A Heap Example



Inserting an Element: `Heapify-up`

1. Insert a new element at $n + 1$
2. Fix the heap order using `Heapify-up($H, n + 1$)`

`Heapify-up(H, i):`

 If $i > 1$ then

 let $j = \text{parent}(i) = \lfloor i/2 \rfloor$

 If $\text{key}[H[i]] < \text{key}[H[j]]$ then


 swap the array entries $H[i]$ and $H[j]$

`Heapify-up(H, j)`


 Endif

 Endif

Heapify-Up Example



Correctness of `Heapify-Up`



- H is **almost a heap with key of $H[i]$ too small** if there is a value $\alpha \geq \text{key}(H[i])$ such that increasing $\text{key}(H[i])$ to α makes H a heap
- **Claim:** The procedure `Heapify-Up`(H, i) fixes the heap property in $O(\log n)$ time, assuming that the array H is almost a heap with the key of $H[i]$ too small.
- **Corollary:** Using `Heapify-Up` we can insert a new element in a heap of n elements in $O(\log n)$ time. (Why?)

Correctness of `Heapify-Up`

- H is almost a heap with key of $H[i]$ too small if there is a value $\alpha \geq \text{key}(H[i])$ such that increasing $\text{key}(H[i])$ to α makes H a heap
- **Claim:** The procedure `Heapify-Up`(H, i) fixes the heap property in $O(\log n)$ time, assuming that the array H is almost a heap with the key of $H[i]$ too small.
- Prove by induction on i :
 - Base case: $i = 1$. $H[1]$ is the root, so if it's too small, then H is already a heap.
 - Inductive step: H is almost a heap with key of $H[i]$ too small. Let $j = \text{parent}(i)$ and β be its key. Swapping the elements at $H[i]$ and $H[j]$ takes $O(1)$ time. After the swap, H is a heap or almost a heap with the key of $H[j]$ too small, since setting its key to β would make H a heap. Finally, by the inductive hypothesis, the recursive call to `Heapify-Up` fixes the heap property.

Deleting an Element: `Heapify-down`

Suppose H has $n + 1$ elements

1. Delete element at $H[i]$ by moving element at $H[n + 1]$ to $H[i]$
2. If element at $H[i]$ is too small, fix heap order using `Heapify-up(H, i)`
3. If element at $H[i]$ is too large, fix heap order using `Heapify-down(H, i)`

`Heapify-down(H, i):`

Let $n = \text{length}(H)$

If $2i > n$ then

 Terminate with H unchanged

Else if $2i < n$ then

 Let $\text{left} = 2i$, and $\text{right} = 2i + 1$

 Let j be the index that minimizes $\text{key}[H[\text{left}]]$ and $\text{key}[H[\text{right}]]$

Else if $2i = n$ then

 Let $j = 2i$


Endif

If $\text{key}[H[j]] < \text{key}[H[i]]$ then


 swap the array entries $H[i]$ and $H[j]$

`Heapify-down(H, j)`

Heapify-down Example



Heapify-down Correctness



- H is **almost a heap with key of $H[i]$ too big** if there is a value $\alpha \leq \text{key}(H[i])$ s.t. decreasing $\text{key}(H[i])$ to α makes H a heap
- **Claim:** The procedure $\text{Heapify-Down}(H, i)$ fixes the heap property in $O(\log n)$ time, assuming that the array H is almost a heap with the key of $H[i]$ too big.
- **Corollary:** Using Heapify-Down we can delete an element from a heap of n elements in $O(\log n)$ time. (Why?)

- H is **almost a heap with key of $H[i]$ too big** if there is a value $\alpha \leq \text{key}(H[i])$ such that decreasing $\text{key}(H[i])$ to α makes H a heap
- **Claim:** The procedure `Heapify-Down`(H, i) fixes the heap property in $O(\log n)$ time, assuming that the array H is almost a heap with the key of $H[i]$ too big.
- Proof by **reverse induction** on i . Suppose H has n elements.
 - Base case: $2i > n$. Then i is a leaf, hence H is a heap.
 - Inductive step: Let j be the child of i with smaller key value and denote its key value β . Swapping the elements at $H[i]$ and $H[j]$ takes $O(1)$ time. The resulting array is a heap or almost a heap with $H[j]$ too big, since setting its key to β makes it a heap. Since $j \geq 2i$, by the inductive hypothesis, the recursive call to `Heapify-Down` fixes the heap property.

In Class Exercise 1

Problem

Naively, we can build a heap out of an arbitrary array using successive calls to HEAPIFY-DOWN, starting at element $\lfloor \text{length}[H]/2 \rfloor$ and going down to 1. If each call to HEAPIFY-DOWN takes $O(\log n)$ time and we have $O(n/2)$ such calls, we can build a heap in $O(n \log n)$ time. Prove that this process is actually faster than $O(n \log n)$ (i.e., provide a *tighter* bound on the process's running time). Starters:

- What is the height of an n -element heap?
- How many nodes are there at height h of an n -element heap?

In Class Exercise 1: continued

What is the height of an n -element heap?

$O(\log n)$ (it's a (nearly) complete binary tree).

In Class Exercise 1: continued

How many nodes are there at height h of an n -element heap?

Key Observation

The number of leaves in a complete binary tree is $\lceil n/2 \rceil$.

Proposition

In an n -element heap, there are $\lceil n/2^{h+1} \rceil$ nodes at height h .

Proof (by induction on h)

Base case: $h = 0$ (the leaves). This is trivially true from the observation above.

Inductive step: Suppose that the claim is true for $h - 1$. Let N_h be the number of nodes at height h in an n -node tree T . Consider T' formed by removing the leaves of T . T' has $n' = n - \lceil n/2 \rceil = \lfloor n/2 \rfloor$ nodes. Nodes at height h in T are at height $h - 1$ in T' (because T' is missing the bottom level of T). Let N'_{h-1} denote the number of nodes at height $h - 1$ in T' .

$$N_h = N'_{h-1} = \lceil n'/2^h \rceil = \lceil \lfloor n/2 \rfloor / 2^h \rceil \leq \lceil (n/2) / 2^h \rceil = \lceil n/2^{h+1} \rceil.$$

In Class Exercise 1: continued

Problem

Naively, we can build a heap out of an arbitrary array using successive calls to HEAPIFY-DOWN, starting at element $\lfloor \text{length}[H]/2 \rfloor$ and going down to 1. If each call to HEAPIFY-DOWN takes $O(\log n)$ time and we have $O(n/2)$ such calls, we can build a heap in $O(n \log n)$ time. Prove that this process is actually faster than $O(n \log n)$ (i.e., provide a *tighter* bound on the process's running time). Starters:

- What is the height of an n -element heap? $O(\log n)$
- How many nodes are there at height h of an n -element heap? $\lceil n/2^{h+1} \rceil$

In Class Exercise 1: Solution

Problem

Naively, we can build a heap out of an arbitrary array using successive calls to HEAPIFY-DOWN, starting at element $\lfloor \text{length}[H]/2 \rfloor$ and going down to 1. If each call to HEAPIFY-DOWN takes $O(\log n)$ time and we have $O(n/2)$ such calls, we can build a heap in $O(n \log n)$ time. Prove that this process is actually faster than $O(n \log n)$ (i.e., provide a *tighter* bound on the process's running time).

Solution

The time required by HEAPIFY-DOWN, when called on a node at height h is $O(h)$. The total cost of building a heap is bounded above by:

$$\sum_{h=0}^{\lfloor \log n \rfloor} \lceil \frac{n}{2^{h+1}} \rceil O(h) = O(n \sum_{h=1}^{\lfloor \log n \rfloor} \frac{h}{2^h}) = O(n)$$

The last step is because (looking up the summation):

$$\sum_{h=0}^{\infty} \frac{h}{2^h} = \frac{1/2}{(1 - 1/2)^2} = 2$$

HeapSort

Sorting with a Priority Queue

Sort

Instance: Nonempty list x_1, x_2, \dots, x_n of integers

Solution: A permutation y_1, y_2, \dots, y_n of x_1, x_2, \dots, x_n such that $y_i \leq y_{i+1}$ for all $1 \leq i < n$

Final Algorithm

- Insert each number in a priority queue H
- Repeatedly find the smallest number in H , output it, and delete it from H

Each insertion and deletion takes $O(\log n)$ time for a total running time of $O(n \log n)$

In Class Exercise 2

Problem

One of your classmates claims that he built an alternative data structure (other than a heap) for representing a priority queue. He claims that, using his new data structure, INSERT, MAX, and EXTRACTMAX all take constant ($O(1)$) time in the worst case. Give a very simple proof that he is mistaken.

Solution

If this were true, we could comparison sort in $O(n)$ time. But we've already proven that this is not possible.

Questions
